

Homework 1

Please scan and upload your assignments on or before February 6, 2020.

- You are encouraged to discuss ideas with each other; but
- you **must acknowledge** your collaborator, and
- you **must compose your own** writeup and/or code independently.
- We **strongly** encourage answers to theory questions in Latex, and answers to coding questions in Python (Jupyter notebooks).
- Maximum score: 50 points.

-
1. **(10 points)** Let $\{x_1, x_2, \dots, x_n\}$ be a set of points in d -dimensional space. Suppose we wish to produce a single point estimate $\mu \in \mathbb{R}^d$ that minimizes the squared-error:

$$\|x_1 - \mu\|_2^2 + \|x_2 - \mu\|_2^2 + \dots + \|x_n - \mu\|_2^2$$

Find a closed form expression for μ and prove that your answer is correct.

2. **(10 points)** Not all norms behave the same; for instance, the ℓ_1 -norm of a vector can be dramatically different from the ℓ_2 -norm, especially in high dimensions. Prove the following norm inequalities for d -dimensional vectors, starting from the definitions provided in class and lecture notes. (Use any algebraic technique/result you like, as long as you cite it.)

- $\|x\|_2 \leq \|x\|_1 \leq \sqrt{d}\|x\|_2$
- $\|x\|_\infty \leq \|x\|_2 \leq \sqrt{d}\|x\|_\infty$
- $\|x\|_\infty \leq \|x\|_1 \leq d\|x\|_\infty$

3. **(10 points)** When we think of a Gaussian distribution (a bell-curve) in 1, 2, or 3 dimensions, the picture that comes to mind is a “blob” with a lot of mass near the origin and exponential decay away from the origin. However, the picture is very different in higher dimensions (and illustrates the counter-intuitive nature of high-dimensional data analysis). In short, we will show that *Gaussian distributions are like soap bubbles*: most of the mass is concentrated near a shell of a given radius, and is empty everywhere else.

- Fix $d = 3$ and generate 10,000 random samples from the standard multi-variate Gaussian distribution defined in \mathbb{R}^d .
- Compute and plot the histogram of Euclidean norms of your samples. Also calculate the average and standard deviation of the norms.
- Increase d on a coarsely spaced log scale all the way up to $d = 1000$ (say $d = 50, 100, 200, 500, 1000$), and repeat parts (a) and (b). Plot the variation of the average and the standard deviation of Euclidean norm of the samples with increasing d .
- What can you conclude from your plot from part (c)?
- Bonus, not for grade.** Mathematically justify your conclusion using a formal proof. You are free to use any familiar laws of probability, algebra, or geometry.

4. **(20 points)** The goal of this problem is to implement a very simple text retrieval system. Given (as input) a database of documents as well as a query document (all provided in an attached .zip file), write a program, in a language of your choice, to find the document in the database that is the best match to the query. Specifically:
- Write a small parser to read each document and convert it into a vector of words.
 - Compute tf-idf values for each word in every document as well as the query.
 - Compute the cosine similarity between tf-idf vectors of each document and the query.
 - Report the document with the maximum similarity value.
5. **(optional)** How much time (in hours) did you spend working on this assignment?